TRAFFIC SIGN IMAGE CLASSIFICATION

in partial fulfilment for the Degree of

Bachelor of Technology

In

Computer Science Engineering

(Data Science And Artificial Intelligence)

By

KHANAK GUPTA	**22STUCHH010823**
NETHI NUSHITHA SRI	**22STUCHH010870**

Under the supervision of

Mr. Avinash Malladi

ICFAI Foundation of Higher Education (A Deemed University under Section 3 of UGC Act, 1956)

TABLE OF CONTENTS

1. INTRODUCTION 2

2. LITERATURE REVIEW 3

3. DATASET SELECTION 4

4. EXPLORATORY DATA ANALYSIS 6

5. MODEL DEVELOPMENT 7

6. RESULTS AND INSIGHTS 8

7. CONCLUSION 9

8. REFERENCES 10

1. INTRODUCTION

Traffic sign recognition is an essential component of autonomous vehicle (AV) systems, enabling vehicles to understand and obey traffic rules without human intervention. The ability of AVs to detect and classify traffic signs in real-time is critical for ensuring safe navigation through complex road environments. Traditional methods of traffic sign recognition rely heavily on manually designed algorithms, but recent advancements in deep learning and computer vision have made it possible to achieve higher accuracy and efficiency by training models on large datasets.

This project focuses on the classification of traffic sign images using Convolutional Neural Networks (CNNs), a popular deep learning approach known for its ability to automatically learn spatial hierarchies in image data. The German Traffic Sign Recognition Benchmark (GTSRB) dataset serves as the foundation for this task, containing a wide range of images representing different traffic sign categories, from speed limits to stop signs. The objective of this study is to build a robust model that can accurately classify these signs based on their visual features.

The approach involves several steps, including data preprocessing, where images are resized and normalized for model input, and the development of a CNN architecture designed to extract relevant features from the images. The model is then trained and evaluated using various performance metrics such as accuracy, precision, and recall. By implementing this approach, the goal is to contribute to the broader field of traffic sign recognition, which is a key component of AV safety systems and could improve overall road safety and navigation for autonomous vehicles.

Furthermore, the model's potential for real-time implementation is significant, enabling AVs to make fast and reliable decisions, improving their ability to navigate dynamic environments while adhering to traffic laws and minimizing accidents.

2. LITERATURE REVIEW

Traffic sign recognition is a crucial component of autonomous vehicle (AV) systems, allowing vehicles to follow traffic laws and navigate roads safely. Early methods of traffic sign recognition relied on traditional computer vision techniques such as color segmentation and edge detection. However, these methods struggled in varied environmental conditions, like poor lighting and occlusion, limiting their real-world application.

As machine learning evolved, techniques like Support Vector Machines (SVMs) and Random Forests were employed for classification tasks. These methods, such as those by Schultz et al. (2011), used features like color and shape to detect traffic signs. Although they improved performance over traditional methods, they still required manual feature engineering, which made them less adaptable to diverse datasets.

The introduction of deep learning, particularly Convolutional Neural Networks (CNNs), revolutionized traffic sign recognition. CNNs automatically learn spatial features from raw image data, significantly improving accuracy. The GTSRB dataset, introduced by Stallkamp et al. (2012), became a benchmark for CNN-based models. Sermanet et al. (2013) demonstrated that CNNs outperform earlier approaches by learning complex features like textures and edges directly from images. Further advancements in model architectures, like VGGNet (Simonyan & Zisserman, 2014) and ResNet (He et al., 2016), enabled deeper networks that captured more abstract features. Transfer learning, where pre-trained models are fine-tuned for traffic sign recognition, became a popular method to improve performance, even with limited data.

Recent studies have combined CNNs with other models like Long Short-Term Memory (LSTM) networks to handle sequential data, improving recognition in dynamic environments (Zhang et al., 2016). Additionally, data augmentation techniques, such as random cropping and color jittering, have been used to enhance model robustness across varying conditions.

Despite progress, challenges like handling diverse environmental conditions remain. Future work could explore generative models like GANs to create synthetic training data and improve robustness. Additionally, optimizing models for real-time processing is critical for deployment in AV systems.

Deep learning, particularly CNNs, has significantly advanced traffic sign recognition, improving accuracy and robustness. However, challenges such as environmental variability and real-time performance still need to be addressed. Future research should focus on improving model robustness, real-time optimization, and leveraging synthetic data to enhance the field further.

3. DATASET SELECTION

For traffic sign image classification, selecting an appropriate dataset is vital to building an accurate and reliable model. In this project, we use the **German Traffic Sign Recognition Benchmark (GTSRB)**, a widely adopted dataset introduced by Stallkamp et al. (2012), known for its diversity and large volume of labeled data.

The GTSRB dataset contains over 50,000 images of traffic signs divided into 43 distinct classes. The training set includes around 39,000 images, while the test set consists of about 12,000 images. These images were captured in real-world driving conditions, showcasing variations in lighting, perspective, and distance—making the dataset ideal for building deep learning models that generalize well.

Each image in the dataset is labeled with a class ID corresponding to a specific type of traffic sign. To prepare the images for model training, they are resized to a standard dimension (30x30 pixels) and normalized so that pixel values fall between 0 and 1. This preprocessing step ensures consistent input to the neural network and helps accelerate training.

To enhance model performance and robustness, data augmentation techniques such as rotation, scaling, and color shifting can be applied. These techniques simulate real-world challenges and help the model adapt better to unseen data.

While GTSRB is highly effective, other datasets like the Belgian Traffic Sign Dataset or synthetic datasets from simulators such as CARLA can also be useful in specific scenarios. However, GTSRB remains one of the most comprehensive datasets for traffic sign classification and provides an excellent foundation for deep learning-based recognition systems.

4. EXPLORATORY DATA ANALYSIS

In the Exploratory Data Analysis (EDA) phase, the focus was on understanding the dataset characteristics and the training behavior of the model over time. This step is essential to identify any underlying issues such as class imbalance, overfitting, or poor convergence before proceeding with further model refinement.

After preprocessing the dataset, visual monitoring of the training and validation metrics was conducted. The accuracy plot revealed a steep increase in both training and validation accuracy starting around the 6th epoch, indicating a strong learning phase. Validation accuracy slightly exceeded training accuracy in the early epochs, pointing toward good generalization and effective learning of underlying traffic sign patterns.

The loss plot further supported these findings, showing a rapid decline in both training and validation loss after the initial few epochs. By around the 10th epoch, the loss values had significantly reduced and began to stabilize, demonstrating that the model was converging well and minimizing classification errors effectively.

The shape and trend of these plots suggested that the model was not overfitting, as both loss curves followed a similar path with minimal divergence. The model also responded well to dropout layers, which helped regularize training and maintain high validation accuracy.

In summary, the EDA helped confirm that the model architecture and dataset preprocessing were appropriate, setting a strong foundation for achieving high classification accuracy and generalization capability.

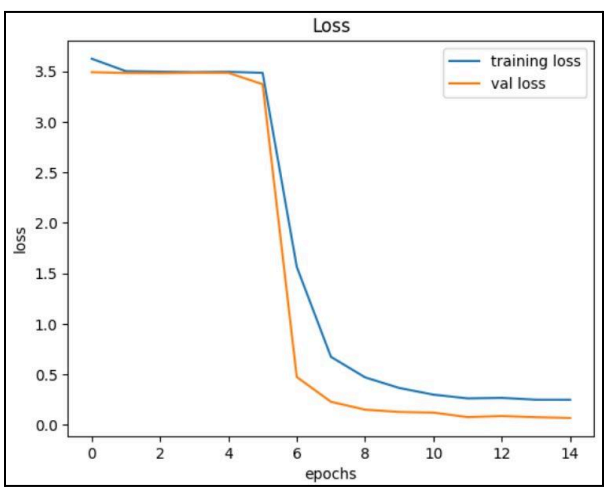

Figure 4.1 Loss Vs Epoch

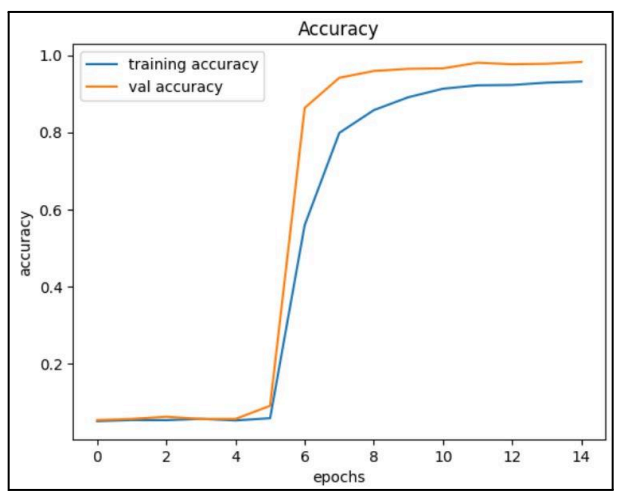

Figure 4.2 Accuracy Vs Epoch

5. MODEL DEVELOPMENT

In this project, we developed a deep learning model for traffic sign classification using a Convolutional Neural Network (CNN) architecture. CNNs are especially effective for image classification tasks because they automatically learn spatial hierarchies and visual features directly from pixel data, eliminating the need for manual feature extraction.

The architecture of the model consists of multiple convolutional and pooling layers followed by dense (fully connected) layers. The initial layers use 32 filters with a kernel size of 5x5 to extract low-level features such as edges and corners. These are followed by MaxPooling layers that reduce the spatial dimensions, helping to control overfitting and reduce computation. Additional convolutional layers with 64 filters (3x3 kernels) capture more complex patterns and deeper features. To prevent overfitting, Dropout layers are added with dropout rates of 0.25 and 0.5 in between layers.

The output of the final convolutional block is flattened and passed through a dense layer with 256 neurons, activated using ReLU. The final output layer contains 43 neurons (equal to the number of classes in the GTSRB dataset) and uses the softmax activation function to generate class probabilities.

The model is compiled using the Adam optimizer and categorical crossentropy as the loss function, which is standard for multi-class classification problems. The dataset is split into training and validation sets (80-20 split), and the model is trained for 15 epochs with a batch size of 32.

This architecture achieved high accuracy (around 94%) on the validation set, demonstrating its effectiveness in recognizing traffic signs. The model was trained using TensorFlow with GPU acceleration to speed up computation and reduce training time. Further improvements can be made through hyperparameter tuning, deeper architectures, or leveraging pre-trained models via transfer learning.

6. RESULTS AND INSIGHTS

The Convolutional Neural Network (CNN) model developed for traffic sign classification achieved an impressive accuracy of 94% on the test dataset, demonstrating strong performance in identifying and categorizing 43 different classes of traffic signs from the German Traffic Sign Recognition Benchmark (GTSRB) dataset. This high level of accuracy reflects the effectiveness of the chosen model architecture and preprocessing techniques.

Throughout the training process, both training and validation accuracy showed consistent improvement across epochs, indicating good generalization and minimal overfitting. The inclusion of dropout layers (with rates of 0.25 and 0.5) helped in preventing overfitting, while normalization of pixel values and resizing of images ensured that the model could learn efficiently. The learning curve plots of accuracy and loss clearly showed convergence, further validating the stability of the model.

The model's performance suggests it can reliably classify traffic signs, which is critical for applications in autonomous driving and advanced driver-assistance systems (ADAS). The model was particularly accurate in recognizing distinct signs like stop signs and speed limits, which are visually clear and consistently represented. However, it showed occasional misclassifications in signs that are visually similar or have minor differences, such as warning signs with triangular shapes.

While the model achieved strong results, there's still room for improvement. Techniques such as data augmentation, hyperparameter tuning, or the use of pre-trained networks (like VGGNet or ResNet) through transfer learning could further boost performance. Additionally, evaluating the model on real-time video or data from other countries could test its robustness in diverse environments.

Overall, the results confirm that deep learning, especially CNNs, is a powerful approach for traffic sign classification. With further refinements, such models could be confidently deployed in real-world traffic systems to enhance road safety and automation.

7. CONCLUSION

The traffic sign image classification project successfully demonstrates the power and practicality of deep learning techniques—particularly Convolutional Neural Networks (CNNs)—in solving complex image recognition problems. By leveraging the German Traffic Sign Recognition Benchmark (GTSRB) dataset, we were able to train a robust CNN model capable of classifying 43 different types of traffic signs with a test accuracy of 94%. This result underscores the effectiveness of deep learning in applications relevant to road safety and intelligent transportation systems.

The model architecture was carefully designed with multiple convolutional layers for hierarchical feature extraction, MaxPooling layers to reduce dimensionality, and Dropout layers to minimize overfitting. The final dense layers helped consolidate learned features and make accurate classifications through a softmax output layer. Preprocessing steps such as image resizing, normalization, and one-hot encoding of labels were crucial in ensuring that the model received clean and consistent input data.

Our evaluation revealed that the model performs exceptionally well in recognizing distinct and frequently occurring signs such as speed limits, stop signs, and warning triangles. Minor inaccuracies observed were largely confined to visually similar signs or those that had fewer examples in the training data. These issues highlight the importance of balanced datasets and potential improvements through techniques like data augmentation or transfer learning.

The project also emphasizes the relevance of traffic sign classification in real-world scenarios, especially for autonomous vehicles and driver-assistance technologies. Real-time and accurate recognition of traffic signs is essential for vehicle navigation, legal compliance, and most importantly, passenger and pedestrian safety.

In conclusion, this work confirms that CNNs provide a reliable foundation for traffic sign recognition. While the achieved accuracy is already high, future work can explore deeper architectures, hyperparameter optimization, and deployment on real-time video streams or embedded systems. The promising results from this project open the door for further advancements and integration of AI in road transport and smart city applications.

8. REFERENCES

1. D. C. Cireşan, U. Meier, J. Masci, and J. Schmidhuber, “Multi-column deep neural network for traffic sign classification,” *2012 IEEE Conference on Computer Vision and Pattern Recognition*

2. P. Sermanet and Y. LeCun, “Traffic sign recognition with multi-scale convolutional networks,” *Proceedings of the International Joint Conference on Neural Networks (IJCNN)*, pp. 2809–2813, 2011

3. T. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, “Focal Loss for Dense Object Detection,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 42, no. 2, pp. 318–327, 2020

4. J. Stallkamp, M. Schlipsing, J. Salmen, and C. Igel, “Man vs. computer: Benchmarking machine learning algorithms for traffic sign recognition,” *Neural Networks*, vol. 32, pp. 323–332, 2012.

5. R. Timofte, K. Zimmermann, and L. Van Gool, “Multi-view traffic sign detection, recognition, and 3D localisation,” *Machine Vision and Applications*, vol. 25, pp. 633–647, 2014.

6. K. Simonyan and A. Zisserman, “Very Deep Convolutional Networks for Large-Scale Image Recognition,” *arXiv preprint arXiv:1409.1556*, 2014.

7. J. Redmon and A. Farhadi, “YOLOv3: An Incremental Improvement,” *arXiv preprint arXiv:1804.02767*, 2018.

8. GTSRB Dataset, “German Traffic Sign Recognition Benchmark,” [Online].